

From Source Availability to Reader Readiness

WikiGraphRAG v18: a research preparation update

4 October 2026

Working addendum | preparation and failure analysis | no natural QA evaluation

Abstract

This update separates source inspection, candidate representation and executable reading as prerequisites for an evidence-reader study. Additional views and literal cover identities now have qualified admission. A versioned interface retains blank blocks and original source projections. After an initial observability failure, a logging-only follow-up observes zero rollback state and tokenizes all authored requests. No answers or natural-question results are produced; complete semantic evidence, reader competence, QA gains and retrieval novelty remain unestablished.

1. Source views and candidate representation

Reviewers inspected additional whole blocks and caption rows against source and CSS records. Their disjoint model-role assignments are not human annotation or an agreement study. Office-rendered views remain qualified inspection aids with disclosed border and browser-fidelity limits. Separate caption recovery clears a missing-view condition, not semantic attachment.

The interface amendment permits empty or whitespace source blocks only. Witness, label and required aspect text remain nonempty. The shape-only successor uses the same assembly and projection helpers and compares each source-only projection, registry and dependency plan with its frozen parent before writing. All original machine bindings retain their declared-metadata status.

Stage	Observed	Meaning
Initial dependency views	39 / 40	Original telemetry failure remains recorded.
Separate caption recovery	1 / 1	Only the previously missing caption is recovered.
Derived dependency views	40 / 40	Combined availability across the two attempts.
Separate form-view recovery	1 / 1	The earlier failed identity-view attempt stays unchanged.
Identity-candidate views	34 / 34	Available views do not automatically admit a card.
New interface shape checks	21 / 21	Representation only; no semantic or author release.

Each row has its own denominator. Derived view availability preserves failed attempts and does not admit semantic support or authors.

Frozen v16 replay preserves ten passes and eleven blank-text failures. All eight exact parent-comparison categories pass for each successor. Forty dependencies, nineteen blank exposures and 326 bundle occurrence entries remain; the last is not a unique-occurrence audit. The shape-only outputs retain null identity cards and populations. Shape validity establishes neither complete evidence nor answerability.

2. From loading failure to authored tokenization

The new application policy retains the kernel hard cap, owned-process-group RSS limit, pressure proxy and reserve, address-space limit, time limits and OOM-event stops. It relaxes the former total-boundary stop (current $\geq$ max): reaching the total-memory boundary and increasing max events are recorded diagnostically. This is a real policy change with shared-cgroup risk, not proof of fit or equivalent isolation.

The first v18 attempt reaches loaded health and identity responses, then fails closed because its required rollback-state log field is unobserved. That is not evidence of nonzero state. Pinned-source diagnosis identifies suppressed runtime log messages. A separately reviewed follow-up changes logging verbosity only, retaining the same assets, authored requests and limits; it observes zero rollback state and tokenizes every authored request. Both attempts confirm cleanup.

Attempt	Last observed gate	Processes	Tokenized	Generation
v18	Health passed; required rollback value unobserved	1	0 / 8	0
v18.1	Health passed; rollback zero; authored tokenization passed	1	8 / 8	0

The logging-only follow-up preserves the earlier failure. Its pass covers loading and these authored prompt tokenizations only.

Observed peak or event count	v18	v18.1
Owned-group RSS (bytes)	4322750464	4321169408
Pressure proxy (bytes)	6724767744	5525028864
Cgroup current (bytes)	8589934592	7896072192
Observed max-event delta	11614	0
Observed OOM-event delta	0	0
Authored input token range	Not measured	907-1315

RSS follows the anchored owned process group. Other peaks and events describe the shared cgroup. Event deltas span first pre-asset through final periodic samples, not after cleanup; these attempts do not isolate model memory or prove reclaimability.

The follow-up also observes rounded runtime buffer reports, which are not exact allocator-byte measurements or proof of a complete generation working set. Server-default generation settings in props and configured future per-completion settings remain unexercised: no generation request or natural prompt was sent. Different shared resource states preclude an efficiency comparison between attempts.

3. Semantic readiness and the next scientific gate

A separate source-only review covers every fixed candidate bundle and retains unknowns for currency identity, exact duration, population, entity, beginning balances and unlabeled row concepts. The review has no typed-record access or private numerical registry join; prior public inventory-candidate metadata exposure is disclosed. Its partial source-visible candidates are not human gold or complete semantic maps.

The initial snapshot treats unresolved codes as descriptions within the fixed source-only projection, not universal reasons that every future machine-assisted question must fail. Separate visible identity candidates are followed by mechanical reconciliation and admission of literal cover identity only. Essential fields retain their inspected source witnesses; optional hidden fields and legal-entity alias inference are excluded. Prior visual review, root literal reading and mechanical reconciliation are distinct roles, not a new visual inspection or agreement study. Author release remains false, and no numerical registry or table entity, period, unit or population binding is admitted. Caption recovery clears the missing-view condition without revising the earlier snapshot.

Separate source-only stages	Count	Boundary
Fixed bundles reviewed	21	No new selection or dependency expansion.
Selected presentation slots	231	Digit-bearing displays and dashes are both retained.
Typed occurrences joined	0	No source-to-machine per-occurrence certification.
Complete semantic maps	0	Unknowns remain explicit.
Complete bundles admitted	0	View availability alone does not clear this gate.
Separate identity candidates	12	Literal visible issuer, report type and end date.
Supported essential fields	36	No optional hidden fields are added.
Literal cover cards admitted	12	No author release or table-level binding follows.

The initial semantic review, later identity candidates and separate literal-card admission are distinct stages. Presentation slots are not typed occurrences. No complete semantic map or pack is admitted.

The next scientific gate is an executable reader on frozen complete evidence, followed by independently authored questions/references and a matched lexical, model-proposal and direct-reader comparison. All arms must see the same acquired evidence; schema validity does not certify question meaning. Natural input fit, missing-evidence controls, complete denominators and residual semantic grading remain necessary. The planned arms receive canonical acquired quantities and therefore test selection and scope, not raw numeric decoding or end-to-end retrieval success.

No retrieval novelty, QA gain, stronger reader, held-out generalization or ACL-ready empirical contribution is established by this checkpoint. The canonical manuscript binds every printed numerical table cell and evidence file hash. Source-bearing artifacts remain external. The PDF uses ReportLab; the editable generic TeX is uncompiled and ACL template compliance is unverified. Earlier manuscripts and failures remain unchanged.