

From Source Occurrences to an Executable Reader

WikiGraphRAG v19: research readiness and retained failures

4 October 2026

Working addendum | preparation and failure analysis | no natural QA evaluation

Abstract

Surviving receipts report a failed reader gate: eight fictional requests complete technically, but only the two single-quantity cases are correct. Source-only projection and initial annotation cover every retained occurrence, while most entries retain unknown relations. Joint scoring has independently reviewed invented controls. After grading, the original native artifacts became unavailable for an unknown reason; surviving public receipts support the reported counts, but a new terminal audit is blocked. This checkpoint establishes neither a stronger reader nor natural QA gains, retrieval novelty or complete raw-artifact reproducibility.

1. Complete projection without semantic promotion

The source-only projection retains fixed bundles, whole contextual blocks, literal source segments and opaque occurrence handles. It excludes normalized values and typed semantic helpers. Exact location and cell membership establish provenance, not visibility or concept, entity, period, unit and population meaning. Entries outside cells remain explicit.

Initial authored schema failures are retained after pre-freeze fixes for metadata aliases and non-text row/cell payloads. A separate failed metadata preflight is also preserved: its versioned correction follows the previously qualified contrast-recovery lineage. This receipt join performs no new source selection or visual inspection.

Retained population or stage	Observed	Boundary
Fixed projected bundles	21	Complete retained bundle population.
Occurrence entries	326	Bundle entries; no deduplication claim.
Nearest table-cell links	228	Cell membership is structural only.
Explicitly outside table cells	98	Retained without a cell claim.
Whole-block bundle links	160	Repeated context is retained.
Unique contextual blocks	141	Fixed view population.
Initial metadata preflight	133 / 141	Failed lineage interpretation retained.
Amended metadata preflight	141 / 141	Receipt join only; no new visual review.

These denominators describe projection and provenance, not complete semantic evidence or answerable questions.

The terminal replay verifies retained artifacts and parent provenance. It is a separate producer-role check, not independent human annotation or a new source/raster review. No deduplication occurred: bundle entries, repeated block links and unique contextual blocks are different populations.

2. Reference-free authored completion

The prediction worker uses pinned inference inputs and fresh processes with fixed decoding, cold prompt caches and continuous resource counters. The recorded workflow seals raw outputs before separate reference-based grading. The original batch has one technical pass, one prelaunch failure and six unattempted rows. Its OSError lacks a retained errno or exact stage.

An owned-loopback fixture supports a possible TIME_WAIT mechanism without proving the original cause. The separate follow-up changes preflight SO_REUSEADDR and stage/errno diagnostics, retaining all inputs, assets and limits. The previous correct answer contributes nothing to its new gate.

Attempt	Planned	Technical pass	Correct	Output limits
v19	8	1	1	0
v19.1	8	8	2	0

Surviving offline-grade receipts report these counts. Both authored gates fail; prior answers are not pooled. Natural prompts remain zero.

Fictional scenario	Proposal reader (B1)	Direct reader (B2)
Single quantity	Answered; correct	Answered; correct
Requested versions	Insufficient; incorrect	Insufficient; incorrect
Ambiguous version	Answered; incorrect	Answered; incorrect
Missing binding	Invalid output	Answered; unsupported

Final actions from the sealed offline grade. These are authored interface controls, not natural questions or a comparative efficacy estimate.

The surviving grade reports technical success but authored-gate failure for the follow-up. Both arms abstain on requested versions and answer both quantities on ambiguity. The missing-binding proposal is invalid; the direct response asserts an unsupported quantity. Neither schema validity nor acquired-source support certifies task correctness.

After root successfully graded the follow-up, original v19 and v19.1 raw native artifacts were found unavailable; the cause is unknown. Public receipts survive. The original v19 terminal audit passed before the loss; the new v19.1 audit is blocked before trace verification. No current raw-retention or full-reproducibility claim is made. These receipt-supported diagnostics establish no arm advantage or natural competence; natural prediction remains blocked.

3. Semantic attachment and scientific limits

Initial source-only annotation covers every retained entry using disjoint assignments across three model roles, including two fresh subroles. The coordinator discloses prior source exposure. This is neither duplicate human annotation nor an agreement study. The frozen snapshot retains local unknowns and an ambiguous presentation; coverage completeness is separate from semantic resolution.

Mechanical checks confirm coverage, schema and witness/page provenance. They do not reconcile the substantive relations or compare them with typed records. No semantic pack, question, reference or author release follows from this snapshot. Unknowns are local to the fixed projection and do not automatically reject every future declared-aspect evaluation.

Initial annotation population	Count	Interpretation
Disposed retained entries	326	Complete initial coverage only.
Reviewed with unknowns	279	One or more source-only relations unresolved.
Reviewed status	47	Initial role disposition; not admitted semantics.
Visible presentation status	325	Source-only role judgment.
Ambiguous presentation	1	Uncertainty is retained.
Witness records	1125	Mechanical provenance checked.
Closed page links	160	Repeated links retained.
Semantic admissions	0	No complete pack or author release.

These initial statuses are not substantive reconciliation, typed comparison, human gold or natural-question evaluation.

The joint scorer requires exact values, complete scopes and acceptable occurrence groups, correlated sufficient citations, and no extra material quantity. Owner and independent invented controls pass. Malformed or failed observed content remains review-required with unknown risk. Pack-supported wrong-target claims may still be incorrect; zero exposed-support errors is not a correctness or safety certificate. Full-study ledger and risk aggregation remain unimplemented.

Next work requires substantive source reconciliation, a separately designed reader correction, and restored durable execution capture before a new frozen evaluation. Matched evidence access and complete planned/admitted/executed denominators remain necessary. No algorithmic novelty, QA gain or ACL-ready contribution is established. Public receipt hashes bind this manuscript. Source annotations remain separate; original native raw artifacts are unavailable. ReportLab produces the inspected PDF; generic TeX is uncompiled and ACL-template compliance is unverified. Earlier receipts and papers remain unchanged.